

The Buffett Quality–Value Ranking

Method, Implementation, and a Point-in-Time Backtest

United States listed equities, 2013–2025

Investa research note • 28 July 2026

Executive summary

This note documents a systematic implementation of Warren Buffett’s investment criteria as a cross-sectional ranking of the entire US listed universe, and reports a point-in-time backtest of the resulting strategy over thirteen years.

1. **The production blend does not beat the market.** Holding the top 20 names by the shipped 60/40 quality–value composite, equally weighted and rebalanced each January, returned **13.6%** a year over 2013–2025 against **14.7%** for the S&P 500 total return index.
 2. **Shifting the blend to 80% quality lifts it to 16.9% a year**, ahead of the S&P 500 by 2.2 points with a smaller maximum drawdown (–23.9% against –27.5%). This is the single change supported by the evidence, and it is supported by a monotone gradient present in the training window, the held-out window and the full sample alike.
 3. **Almost everything else that was tested was noise.** A sweep of 900 configurations — concentration, market-capitalisation floors, momentum overlays, two-stage screens — produced a set of training-window winners whose median held-out return (12.0%) was *below* that of the grid as a whole (12.4%). Tuning made things worse out of sample.
 4. **No configuration beat the NASDAQ-100 (19.7% a year).** The screen never held the mega-cap technology names that drove that index: across 2020–2025 a Magnificent Seven stock entered the top 20 in one year out of six.
 5. **The results are flattered by survivorship bias** and cover thirteen years rather than twenty, because SEC XBRL fundamentals do not reach further back. Both limits are quantified in §6.
-

Contents

Executive summary	1
1 The Buffett method, and what it demands of a screen	4
1.1 Quality first, price second	4
1.2 Owner earnings, not accounting profit	4
1.3 Durability, not the latest year	4
1.4 Circle of competence, expressed as sector models	4
1.5 Margin of safety	5
2 Design principles	5
3 Implementation	5
3.1 Universe	6
3.2 Fundamentals: SEC EDGAR XBRL	6
3.3 Quality: five pillars	6
3.4 Gates	6
3.5 Sector models	7
3.6 Value	8
3.7 The composite	8
4 Backtest methodology	9
4.1 Point-in-time fundamentals	9
4.2 Prices in the money of the day	9
4.3 Portfolio construction	9
5 Results	9
5.1 The production blend	10
5.2 The parameter search, and its main negative result	10
5.3 The one change that survives	11
5.4 Industry concentration	11
5.5 The Magnificent Seven	12
6 Limitations	13
6.1 Survivorship bias	13

6.2	Thirteen years, not twenty	13
6.3	A relaxed history gate	13
6.4	What the backtest does not model	14
7	Conclusion	14
A	Reproducing these results	14
B	Current portfolio	15

1 The Buffett method, and what it demands of a screen

Buffett's public writing is remarkably consistent about what he is looking for, and remarkably unlike what most quantitative value screens actually measure. A faithful implementation has to take five ideas seriously.

1.1 Quality first, price second

The ordering is not a stylistic preference; it is the whole argument. In the 1989 letter to shareholders Buffett described the lesson of his own worst purchases: "It's far better to buy a wonderful company at a fair price than a fair company at a wonderful price." A screen that ranks on discount-to-value inverts this. It will reliably put at the top the companies the market expects to shrink, because a business in permanent decline looks cheap against any model that assumes it will not decline.

This has a concrete architectural consequence, described in §3.4: quality tests run as *gates* before any price is considered, and a company that fails one is removed from the ranking entirely rather than scored badly and buried.

1.2 Owner earnings, not accounting profit

Buffett's 1986 appendix defines owner earnings as reported earnings plus depreciation and amortisation, less the capital expenditure required to maintain competitive position. The final term is the difficulty: no filer discloses maintenance capital expenditure separately from growth capital expenditure. The implementation substitutes total capital expenditure:

$$\text{owner earnings} = \text{operating cash flow} - \text{capital expenditure}$$

which understates the earnings of a company investing heavily in growth. That is a deliberate direction to err in — a quality screen should be conservative about companies whose cash generation depends on assumptions.

1.3 Durability, not the latest year

A 30% return on equity in one year is noise. Fifteen per cent in each of twelve years is evidence of something structural. Consequently almost nothing in the metric set is measured on the most recent year: profitability, predictability and growth are medians, dispersions and long-window compound growth rates computed over a ten-year window. The exception is genuine balance-sheet *state* — what the company owes today — which is read from the latest period, because that is what it is.

This requirement is also what makes the data problem hard, and it is the binding constraint on how far back the backtest can reach (§6.2).

1.4 Circle of competence, expressed as sector models

Buffett's circle-of-competence discipline is usually read as advice to individuals. For a screen it translates into something sharper: a valuation model applied outside its domain produces a confident, meaningless number. A discounted-cash-flow valuation of a bank is not a poor estimate; it is an answer to the wrong question, because for a deposit-taking institution liabilities

are the raw material of the business rather than a financing choice, and “free cash flow” is not a coherent quantity.

The implementation routes every company to one of four models on its SEC Standard Industrial Classification code (§3.5).

1.5 Margin of safety

Graham’s principle, which Buffett calls the three most important words in investing, is that the estimate of value must exceed the price by enough to absorb the estimate being wrong. This is implemented not as a haircut to a point estimate but as a rank on the 25th percentile of a Monte Carlo distribution of intrinsic value (§3.6), on the reasoning that the mean of a valuation distribution is not a price anyone should pay.

2 Design principles

Eight principles govern the implementation. They are stated here because several of them constrain the code in ways that are not obvious from the outputs.

P1 — Gates before scoring.

Solvency and persistent-loss tests remove a company from the ranked list and report it separately with a reason, rather than contributing a poor score.

P2 — Conservative estimates over central ones.

Value is ranked on a low percentile of the valuation distribution; dispersion is penalised separately.

P3 — Reuse the existing engine.

Ratios come from the same `financial_ratios` module the rest of the application uses, so the ranking and the stock-detail view can never disagree about a company’s ROE.

P4 — Durability over level.

Medians and dispersions over a ten-year window, not point values.

P5 — Explainability.

Every score decomposes into pillars and metrics, and every exclusion carries a machine-readable reason.

P6 — Compare like with like.

All percentiles are computed *within* a valuation model. A bank’s equity-to-assets ratio is meaningful against other banks and meaningless against software companies.

P7 — Missing data must never look good.

Every metric returns “unknown” rather than a default; unknowns flow into a confidence factor that can only demote. A gate fires only on a known violation.

P8 — Point-in-time honesty.

Facts are stored keyed by filing accession, so the originally reported number remains available after a restatement. This is what makes §5 a backtest rather than a hindsight exercise.

3 Implementation

3.1 Universe

The investable universe is built from the NASDAQ Trader symbol directories (`nasdaqlisted.txt` and `otherlisted.txt`), which together list every symbol on every US exchange. Filtering removes exchange-traded funds, warrants, rights, units, preferred series, depositary receipts and test issues — roughly 13,000 raw rows reduce to about 5,600 common-stock listings. Each is matched to an SEC Central Index Key so fundamentals can be retrieved. Multiple share classes of one filer (BRK-A/BRK-B, GOOG/GOOGL) collapse to a single company, since they share one set of financial statements.

3.2 Fundamentals: SEC EDGAR XBRL

The data source is the SEC’s XBRL company-facts archive rather than a commercial vendor, for a specific reason: vendor feeds available here return four to five annual periods per company, and §1.3 requires ten. EDGAR carries roughly seventeen annual periods per filer at no cost.

Three implementation details matter for correctness:

- **Only annual facts from annual reports** are retained. Quarterlies would let a company’s fiscal calendar leak into cross-sectional comparisons. A “duration” fact counts as annual if it spans 300 to 400 days, which accommodates the 52- and 53-week retail calendars.
- **Concepts resolve through fallback chains, period by period.** XBRL tag usage varies by company, by year and above all by industry. Apple’s revenue arrives under `Revenues` for older years and `RevenueFromContractWithCustomer...` for newer ones; resolving per period rather than per company recovers one continuous series across the accounting-standard change.
- **Every (CIK, tag, period, accession) is its own row.** A restatement never overwrites the originally filed value. The resolver defaults to the most recent filing, but the originals remain — which is precisely the mechanism §5.1 depends on.

3.3 Quality: five pillars

Each metric is converted to a percentile within its valuation model, winsorised at the 1st and 99th centiles so that a single absurd outlier — a 10,000% return on equity computed from a near-zero equity base — cannot compress every other company into a narrow band. A pillar score is the mean of whichever of its metrics resolved; the quality score is the weighted mean of the pillars, renormalised over those present.

Two metrics deserve comment. **Return on invested capital** nets cash off the capital base, because a large idle balance would otherwise depress the measured return on the operating business the moat question is actually about. **Incremental return on incremental capital** — the change in pre-tax earnings divided by the change in equity across the window — is the sharpest available test of whether management creates value with retained earnings. It is preferred to Buffett’s own “one dollar retained produces one dollar of market value” formulation because it requires no price history and therefore cannot be contaminated by look-ahead bias.

3.4 Gates

Gates encode P1 and P7 together: a company is excluded for something its filings *show*, never for something they omit. Every test is written so that a missing value cannot trigger it.

Pillar	Weight	Metrics (generic model)
Returns on capital	30%	Median ROE, median ROIC, median ROA, median gross margin, share of years with ROE above 15%
Financial strength	20%	Debt-to-equity, interest coverage, current ratio, net debt to owner earnings
Predictability	20%	Dispersion of ROE, of revenue growth, of FCF margin; count of negative owner-earnings years
Growth	15%	Revenue CAGR, owner-earnings CAGR, book-value-per-share CAGR
Capital allocation	15%	Share-count CAGR (falling is good), incremental return on incremental capital

Table 1: Quality pillars and weights. Returns on capital carries the most weight because a durable high return on capital *is* the moat; growth carries the least because growth without returns destroys value.

Gate	Threshold	Applies to
Insufficient history	< 8 annual periods	all (relaxed to 5 in the backtest, \$6.3)
Interest not covered	coverage < 2.0	non-financials that report debt
Debt not serviceable	net debt / owner earnings > 6.0	non-financials
Persistent cash burn	> 2 years of negative owner earnings	non-financials
Unprofitable	median ROE < 0	all except REITs
Undercapitalised	equity/assets < 4%	banks, insurers
Persistent losses	> 2 loss years	banks, insurers
Excessive leverage	debt/FFO > 15	REITs

Table 2: Quality gates. Roughly four fifths of the listed universe fails at least one.

The most instructive gate is the one that was *removed*. Debt-to-equity above 2.0 was the obvious solvency test and it was wrong: sustained buybacks shrink book equity, so the ratio explodes for precisely the highest-quality compounders. At that threshold it excluded 253 companies on that basis alone, among them Apple — which carried a debt-to-equity ratio of 3.9 alongside \$99bn of annual owner earnings against \$55bn of net debt — as well as ADP, Amgen and American Express. Book equity is an accounting residual; the question that matters is whether the business generates enough cash to retire what it owes. Debt-to-equity survives as a *scoring* input inside the financial strength pillar, where a poor reading costs percentile points rather than removing a company from consideration.

3.5 Sector models

Model	SIC range	Basis of measurement
Bank	6020–6036, 6060–6062, 6080–6082, 6110–6120	Residual income; equity/assets, net interest margin, provisioning, efficiency ratio
Insurer	6311–6411	Residual income, without lending metrics
REIT	6798	Funds from operations (NAREIT definition)
Generic	everything else	Free cash flow and the standard pillars

Table 3: Valuation model routing.

SIC is used in preference to a vendor sector string because the distinctions that matter here are exactly the ones vendor classifications blur. A “Financial Services” bucket mixes deposit-taking

banks, which require a residual-income model, with asset managers and exchanges, which have ordinary free-cash-flow economics. Asset managers and brokers (SIC 6199, 6211) are therefore deliberately *not* treated as financials here.

Insurers were separated from banks after measurement showed that 29% of everything routed to the bank model had no deposits, no loan book and no net interest income. Scoring those companies on lending metrics would have demoted every one of them for lacking data they will never report — a data gap masquerading as poor quality, which is exactly what P7 exists to prevent.

For REITs, funds from operations replaces net income ($\text{FFO} = \text{net income} + \text{real-estate depreciation} - \text{gains on property sales}$) because GAAP depreciation is charged on assets that in practice hold or gain value. Share-count growth carries unusual weight in the REIT model: REITs fund themselves by issuing equity, and growing FFO while growing the share count faster is not growth.

3.6 Value

The value score blends a discounted-cash-flow margin of safety with classic multiples, all ranked as percentiles within the valuation model.

Component	Weight	Direction
Margin of safety (P25 of Monte Carlo)	35%	higher is better
Earnings yield	20%	higher is better
Free-cash-flow yield	20%	higher is better
EV/EBIT	10%	lower is better
Price/book	10%	lower is better
Price/sales	5%	lower is better

Table 4: Value components. The multiples carry more weight in aggregate than the discounted cash flow because they rest on reported figures rather than on a projected growth rate.

Two refinements implement P2. First, the intrinsic value taken forward is the 25th percentile of the Monte Carlo distribution rather than the point estimate. Second, a *dispersion penalty* shrinks the margin of safety in proportion to the width of that distribution, computed as $(P_{90} - P_{10})/P_{50}$: a distribution twice as wide as its own midpoint has its margin roughly halved. This is what stops a speculative business with a huge but meaningless apparent discount from topping the value ranking. Where the discounted-cash-flow and Graham models disagree by more than a factor of two, the lower of the two is taken rather than the average, on the view that such disagreement is itself evidence that neither should be trusted at face value.

Negative multiples are treated as missing rather than as cheap. A loss-making company has no meaningful price-to-earnings ratio and a company with negative book value has no meaningful price-to-book ratio; ranking them as “cheapest” would be an arithmetic accident.

3.7 The composite

$$\text{score} = (w_Q \cdot \text{quality} + (1 - w_Q) \cdot \text{value}) \times \text{confidence}$$

with $w_Q = 0.60$ as shipped. The confidence factor maps data coverage — the fraction of the model’s headline metrics that resolved — onto the interval $[0.5, 1.0]$. It is capped at 1.0 so that it can only ever demote: full coverage earns no bonus, it is merely the absence of a penalty. The

floor of 0.5 stops a thinly reported company from collapsing to zero and disappearing silently instead of ranking low and visibly.

Where no value score is available the composite is quality alone, so the ranking degrades gracefully when market data is missing rather than failing.

4 Backtest methodology

4.1 Point-in-time fundamentals

Every fact in the store carries the date it was filed. The backtest opens an as-of context that restricts all reads to facts filed on or before the rebalance date, so the January 2015 ranking sees only what had actually been filed by 31 December 2014 — the originally reported numbers rather than subsequent restatements, and no fiscal year that had not yet been reported.

The mechanism was verified directly: in the 2015 ranking the latest fiscal period visible for any company ends 1 November 2014. Apple’s median return on equity as of that date reads 29.6%, against 119% on today’s data — the difference being a decade of buybacks shrinking the equity base. A backtest using current data would have ranked the 2015 universe on a number that did not exist in 2015.

4.2 Prices in the money of the day

Yahoo’s closing prices are retroactively adjusted for splits. Multiplying such a price by the share count a company reported at the time understates its market capitalisation by the cumulative split factor — fourfold for Apple, fortyfold for NVIDIA. The backtest therefore reconstructs the price actually quoted by multiplying the adjusted close by the product of all splits occurring after that date.

Verified against known closes: Apple on 31 December 2014 reconstructs to \$110.38 against an actual close of \$110.38; NVIDIA to \$20.05 against \$20.03; Simon Property Group to \$182.11 against \$182.03. Share counts come from the latest EDGAR filing available at the as-of date, and as-reported earnings per share are used unmodified, so price, shares and per-share figures are all expressed on the same basis. Returns are measured separately from the dividend-and-split-adjusted series, so they are total returns.

4.3 Portfolio construction

At the first rebalance of each year the top N names by composite score are bought in equal weight and held for the calendar year, with weights left to drift. A position whose price series ends mid-year — delisting or acquisition — is frozen at its last quote and carried as cash for the remainder of the year. Companies without a quote on the rebalance date are dropped rather than ranked on quality alone, since they could not have been bought.

5 Results




Figure 1: Growth of one dollar, 2013–2025, monthly. The strategies are the top 20 names by composite score, equally weighted, rebalanced each January. Benchmarks are total-return series, so dividends are included on both sides.

	Total	CAGR	Volatility	Max drawdown	Return/vol
Ranking 80/20 (top 20)	667.6%	16.9%	17.2%	−23.9%	0.98
Ranking 80/20, sector-capped	613.6%	16.2%	17.1%	−24.3%	0.95
Ranking 60/40 (production)	431.1%	13.6%	17.7%	−27.5%	0.77
S&P 500 (total return)	499.8%	14.7%	14.2%	−23.9%	1.04
NASDAQ-100 (total return)	952.2%	19.7%	17.2%	−32.6%	1.15

Table 5: Summary statistics, 2013–2025. Volatility is annualised from monthly returns; drawdown is measured on the monthly equity curve.

5.1 The production blend

The shipped configuration underperforms. At 60/40 the strategy returned 13.6% a year against the S&P 500's 14.7%, with higher volatility and a deeper drawdown. It beat the index in 7 of 13 years.

The pattern of those wins is coherent rather than random. The screen won 2015, 2016, 2018 and the relative battle in 2022 — the years when growth stalled — and lost heavily in 2019, 2020 and 2024, when mega-cap growth ran. That is what a quality-and-value screen does in a growth decade.

5.2 The parameter search, and its main negative result

Nine hundred configurations were evaluated: quality weight from 0 to 1, concentration from 5 to 50 names, market-capitalisation floors from zero to \$10bn, and 12-1 month momentum both as a filter and as a blended score component. Each was scored on a training window (2013–2019) and re-measured on a held-out window (2020–2025) it had played no part in selecting.

The top 15 configurations by training-window return had a median held-out return of 12.0%, against 12.4% for the grid as a whole. Selecting on the training window produced worse out-of-sample results than choosing at random from the same grid.

Individual configurations swung violently between the two windows: one that returned 24.3% in training delivered 4.3% out of it. Momentum is the clearest illustration: weighting it at 0.5 was the *worst* setting in training (12.4% median) and the *best* in the held-out window (14.6%). That is a regime effect, not an edge, and any conclusion drawn from either window alone would have been wrong.

5.3 The one change that survives




Figure 2: Median compound return across the entire grid, by weight on quality. Each point is the median over every other setting — concentration, market-cap floor, momentum — so it measures the effect of the weight alone rather than of one favourable combination.

The gradient is monotone from 0 to 0.8 in all three windows, and only the quality-only extreme (1.0) collapses out of sample, from a median of 13.9% at 0.8 to 10.0%. A single-parameter change supported by a consistent slope is far weaker evidence of overfitting than the argmax of a noisy 900-cell grid, which is why it is the only change recommended here.

At 80/20 the strategy returns 16.9% a year, beating the S&P 500 by 2.2 points with a marginally smaller drawdown. It improves on the 60/40 blend in nine of thirteen individual years. Concentration is a stable plateau from 10 to 20 names (top 10 returns 18.4%); five names is fragile, returning 24.3% in training and 4.3% in the held-out window.

5.4 Industry concentration

The ranking scores companies one at a time and has no representation of correlation, so it will fill a fifth of the book with companies sharing a single risk factor. In the current live ranking, four of the top twenty are mortgage insurers (SIC 6351) and two more are homebuilders (SIC 1531): 30% of the portfolio on one macroeconomic exposure.

A cap on names per industry was implemented and tested.




Figure 3: Calendar-year returns, 80/20 ranking against the S&P 500 total return index. The strategy beat the index in nine of thirteen years, and the excess is not concentrated in any single year.

Rule (80/20, top 20)	CAGR	Return/vol	Max DD	Train	Held-out
No cap	16.9%	0.98	−23.9%	20.0%	13.1%
Max 3 per 2-digit SIC	16.2%	0.95	−24.3%	19.1%	12.7%
Max 3 per 3-digit SIC	16.1%	0.92	−25.4%	19.8%	11.6%
Max 2 per 3-digit SIC	15.5%	0.89	−27.3%	18.6%	11.8%
Max 1 per 3-digit SIC	15.3%	0.85	−24.5%	18.7%	11.3%

Table 6: The cost of capping industry concentration.

Capping cost return and did not reduce drawdown. Tighter caps were monotonically worse on every measure — the concentrated industries happened to contain the better picks. The caveat cuts the other way: the sample contains no housing crisis, which is precisely the event the cap insures against. A backtest cannot price a tail that did not occur in it.

Note also what a SIC cap cannot see. Mortgage insurance (6351) and homebuilding (1531) are different industries carrying the same housing exposure, and no classification-based cap will group them. Controlling that requires an explicit factor bucket.

5.5 The Magnificent Seven

Across 2020–2025, a Magnificent Seven constituent entered the top 20 in one year: 2023, when Meta ranked 7th and Alphabet 15th.

Their quality scores were strong throughout — Microsoft 67–74, Alphabet 69–74, all around the top decile — and their value scores were punishing: NVIDIA scored 7 in 2021, 13 in 2023 and 4 in 2024. The value half of the composite is what excluded them, which is visible in the weighting itself: Microsoft ranked 336th under 60/40 in 2021 and 93rd under 80/20. The same company, moved 240 places by the weight alone.

For scale, an equal-weight Magnificent Seven portfolio rebalanced annually returned 43.1% a year over 2020–2025, with 32% volatility and a −47% drawdown. No quality-and-value screen was going to hold those names at those multiples, and that single omission accounts for most of

Year	AAPL	MSFT	GOOG	AMZN	NVDA	META	TSLA
2020	261	119	310	731	374	148	—
2021	771	93	222	727	416	171	—
2022	586	78	189	564	337	58	—
2023	559	47	15	403	469	7	—
2024	634	89	228	—	693	219	—
2025	965	79	223	985	34	108	—

Table 7: Rank under the 80/20 blend, out of roughly 1,100–1,250 ranked companies. A dash means the company failed a quality gate that year. Tesla failed one in every year of the sample.

the gap to the NASDAQ-100.

6 Limitations

6.1 Survivorship bias

The universe is constructed from current listing files, so companies delisted, acquired or bankrupted between 2013 and 2026 were never candidates in any year. This flatters the strategy, and it flatters it asymmetrically: the index benchmarks are survivorship-free. Acquisitions typically complete at a premium and bankruptcies at a total loss, and neither appears here. Every strategy figure in §5 should be read as an upper bound. Correcting it requires a historical constituent database, which is not available from free sources.

6.2 Thirteen years, not twenty

XBRL filing became mandatory for large filers only in 2009, and a 10-K carries two years of comparatives. The store therefore holds fundamentals for 32 companies in fiscal 2005 and 298 in fiscal 2007.

As of	Companies with any facts	With ≥ 5 annual periods
31 Dec 2011	865	26
31 Dec 2012	1,966	372
31 Dec 2013	2,042	862
31 Dec 2014	2,137	1,607

Table 8: Point-in-time data availability. 2013 is the earliest rebalance year with a usable cross-section.

6.3 A relaxed history gate

The production gate requires eight annual periods. At the end of 2014 that threshold left 68 eligible companies out of roughly 4,000, which would have made the early years meaningless. The backtest therefore requires five — the same floor the metrics module uses before it will compute a consistency measure at all. Later years are effectively unaffected; the early years use a shorter durability window than production does, which if anything understates the quality signal.

6.4 What the backtest does not model

Transaction costs, bid-ask spreads, market impact and taxes are all excluded. For an annually rebalanced 20-name book the turnover cost is small but not zero. Some holdings are genuinely small companies where the assumption of frictionless execution at a month-end close is optimistic — and those are the same names that survivorship bias most affects.

7 Conclusion

The ranking implements Buffett's criteria faithfully enough to be worth taking seriously, and the exercise produced one durable result and one cautionary one.

The durable result is that **the quality half of the composite carries the signal and the value half dilutes it**. Raising the quality weight from 0.60 to 0.80 moves the strategy from 1.1 points behind the S&P 500 to 2.2 points ahead, and the evidence for it is a monotone gradient present in every window tested rather than the peak of a noisy search.

The cautionary result is that **almost everything else tested was overfitting**. A grid large enough to contain an impressive answer will always contain one; here the training-window winners were measurably worse than average out of sample. Any future change to this ranking should be held to the standard the quality weight met: a consistent effect across independent windows, not a high score in one.

Two open questions follow. The screen has no mechanism to pay up for durability — a company can be the highest-quality business in the universe and still be ranked out on price — and whether that is a flaw or the entire point is a judgement about investment philosophy that a backtest cannot settle. Second, the value component deserves direct investigation: a 40% weight on a component that subtracts return is worth understanding before its scores are trusted elsewhere in the application.

A Reproducing these results

File	Role
src/edgar_provider.py	Fact store; the as_of() context implementing P8
src/buffett_metrics.py	Per-company metrics, all four sector models
src/buffett_rank.py	Gates, pillar percentiles, the composite
src/buffett_value.py	Margin of safety, dispersion penalty, multiples
src/buffett_pipeline.py	End-to-end production run
scripts/buffett_backtest.py	Point-in-time backtest
scripts/buffett_strategy_search.py	Parameter sweep and held-out evaluation
scripts/buffett_live_picks.py	Sized buy list from the latest run

```
python scripts/buffett_backtest.py --start 2013 --end 2025 --top 20
python scripts/buffett_strategy_search.py --start 2013 --end 2025 --split 2019
python scripts/buffett_live_picks.py --capital 200000 --quality 0.8 \
    --max-per-sector 3 --sector-digits 2
```

Per-year rankings are cached, so re-running with different parameters takes seconds rather than minutes.

B Current portfolio

Top 20 under the 80/20 blend with a cap of three names per 2-digit SIC group, from the ranking run of 28 July 2026 (5,539 filers, 1,244 ranked), sized for \$200,000 of capital.

Symbol	Industry	Price	Shares	Cost	Weight
MTG	Mortgage insurance	30.00	333	9,990	5.00%
RDN	Mortgage insurance	39.33	254	9,990	4.99%
ESNT	Mortgage insurance	66.59	150	9,989	4.99%
UNTY	Banks	58.59	170	9,959	4.98%
INVA	Pharma	21.68	461	9,994	5.00%
LULU	Apparel	117.61	85	9,997	5.00%
TSBK	Banks	44.15	226	9,978	4.99%
PKBK	Banks	33.37	299	9,978	4.99%
ADBE	Software	237.72	42	9,984	4.99%
GOOG	Software	319.09	31	9,892	4.95%
ACN	Software & services	154.20	64	9,869	4.93%
PHM	Homebuilders	129.29	77	9,956	4.98%
G	Consulting	32.35	309	9,996	5.00%
XPEL	Fabricated metal	43.66	229	9,999	5.00%
MKTX	Brokers & exchanges	118.12	84	9,922	4.96%
NVR	Homebuilders	6,478.96	1	6,479	3.24%
CPRT	Auto services	29.49	339	9,997	5.00%
ESRT	REITs	5.62	1,779	9,998	5.00%
AGM	Non-bank credit	211.58	47	9,944	4.97%
RMD	Instruments	203.10	49	9,952	4.98%
Invested				195,865	97.9%

Residual concentration after capping: 15.0% mortgage insurance, 15.0% banks, 14.9% software and services, 8.2% homebuilders. Mortgage insurance and homebuilding together remain a 23% exposure to US housing credit.

NVR at \$6,479 a share cannot fill a 5% slot below roughly \$130,000 of capital; fractional-share execution resolves this and puts the residual \$4,135 of cash to work.

This document reports the output of a quantitative model on historical data. It is not investment advice, and past results — particularly backtested results carrying the biases described in §6 — do not indicate future returns.